

# Towards Closing the Energy Gap Between HOG and CNN Features for Embedded Vision

Amr Suleiman\*, Yu-Hsin Chen\*, Joel Emer, Vivienne Sze

Massachusetts Institute of Technology



Contact Info

email: [sze@mit.edu](mailto:sze@mit.edu)

website: [www.rle.mit.edu/eems](http://www.rle.mit.edu/eems)



Amr Suleiman

Yu-Hsin Chen

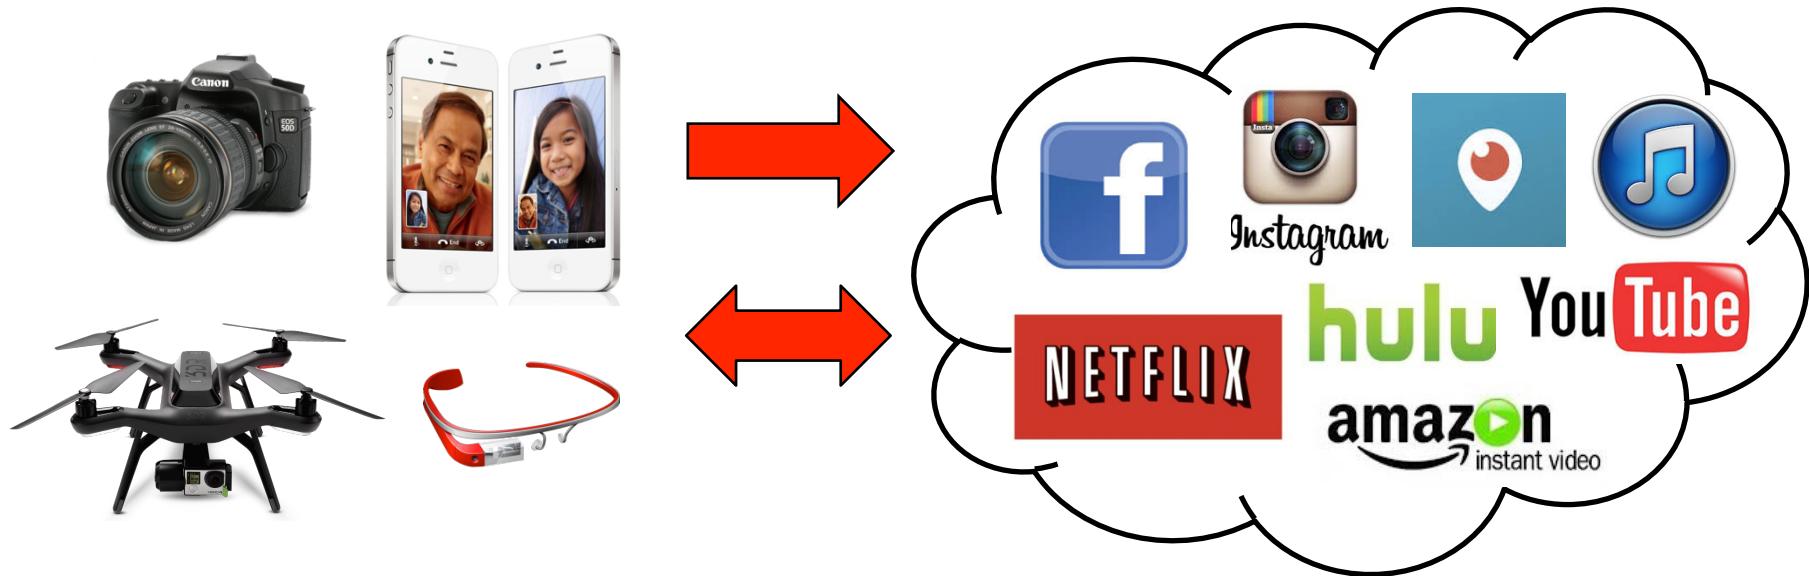


# Video is the Biggest Big Data

Over 70% of today's Internet traffic is video

Over 300 hours of video uploaded to YouTube **every minute**

Over 500 million hours of video surveillance collected **every day**



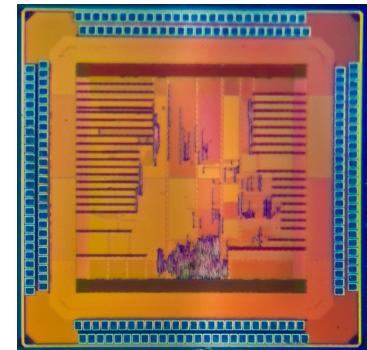
*Energy limited due  
to battery capacity*

*Power limited due  
to heat dissipation*

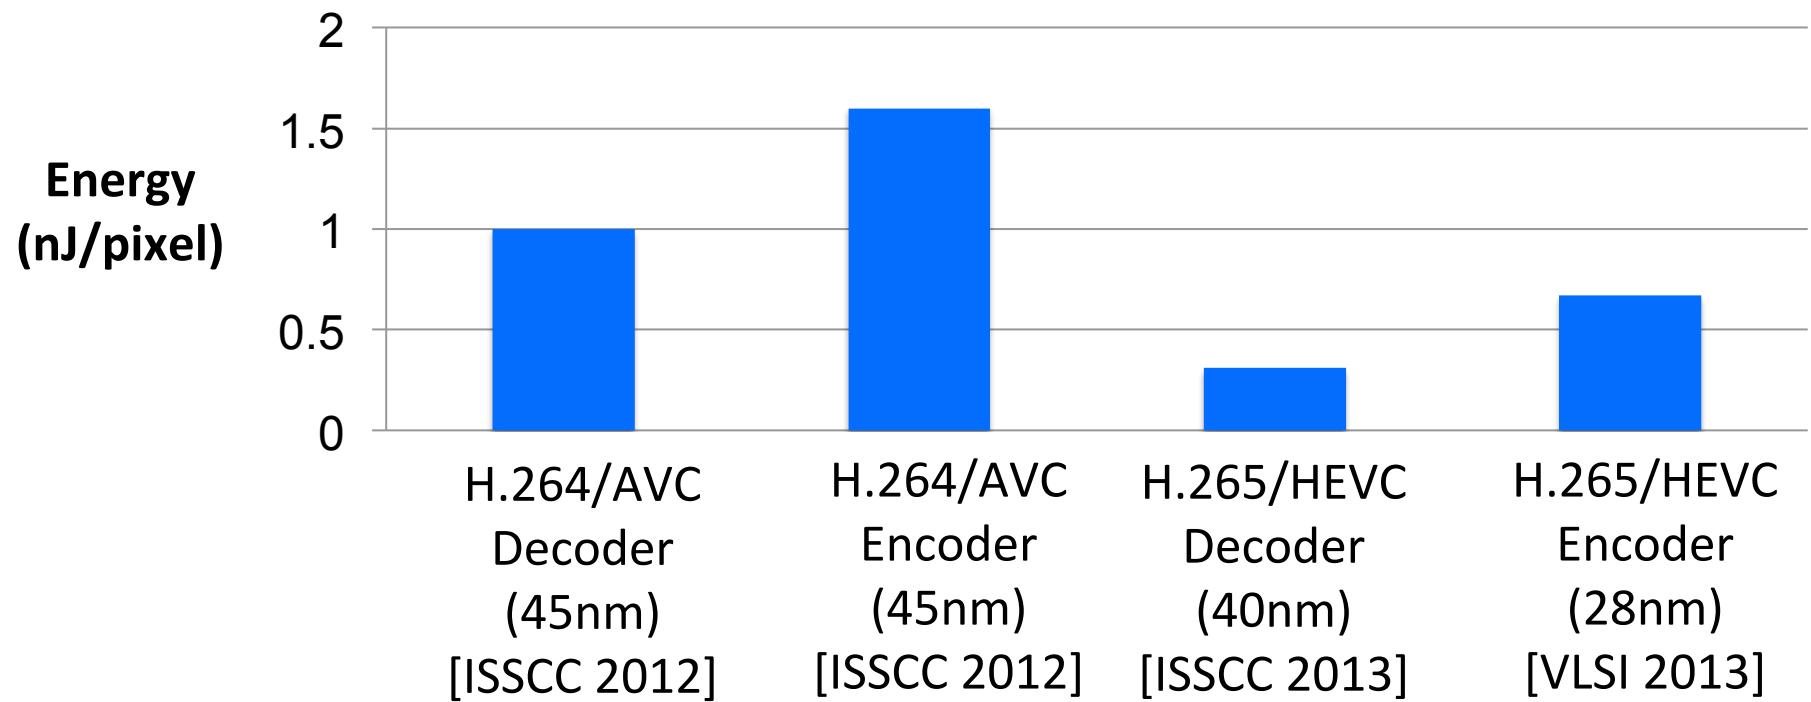
Need energy-efficient pixel processing!

# Typical Constraints on Video Compression

- **Area cost:** Memory Size of 100-500kB, ~1000kgates
- **Power budget:** < 1W for smartphones
- **Throughput:** Real-time 30 fps
- **Energy:** ~1nJ/pixel

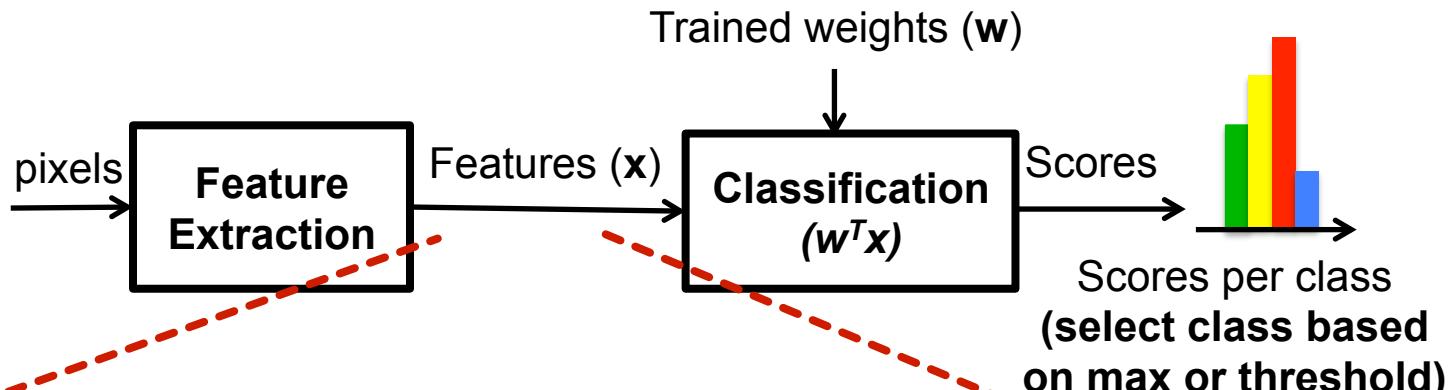


[ISSCC 2014]

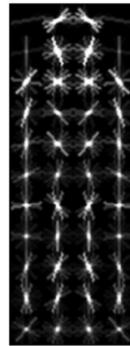
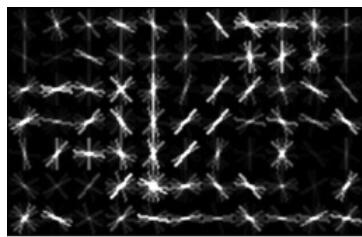


# Object Detection/Classification Pipeline

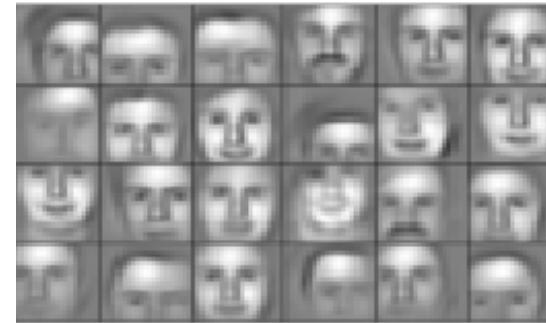
Image



Handcrafted Features  
(e.g. HOG)



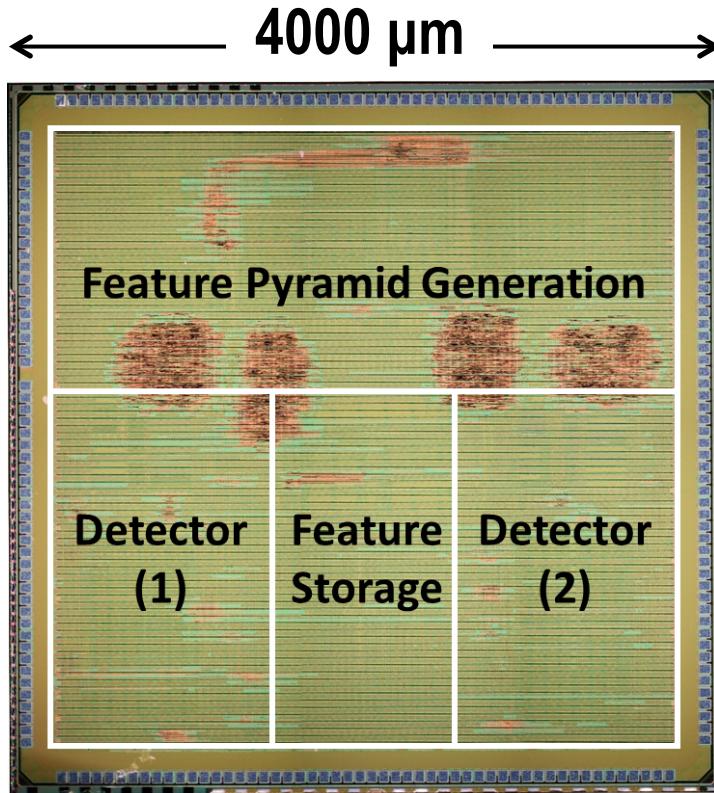
Learned Features  
(e.g. CNN)



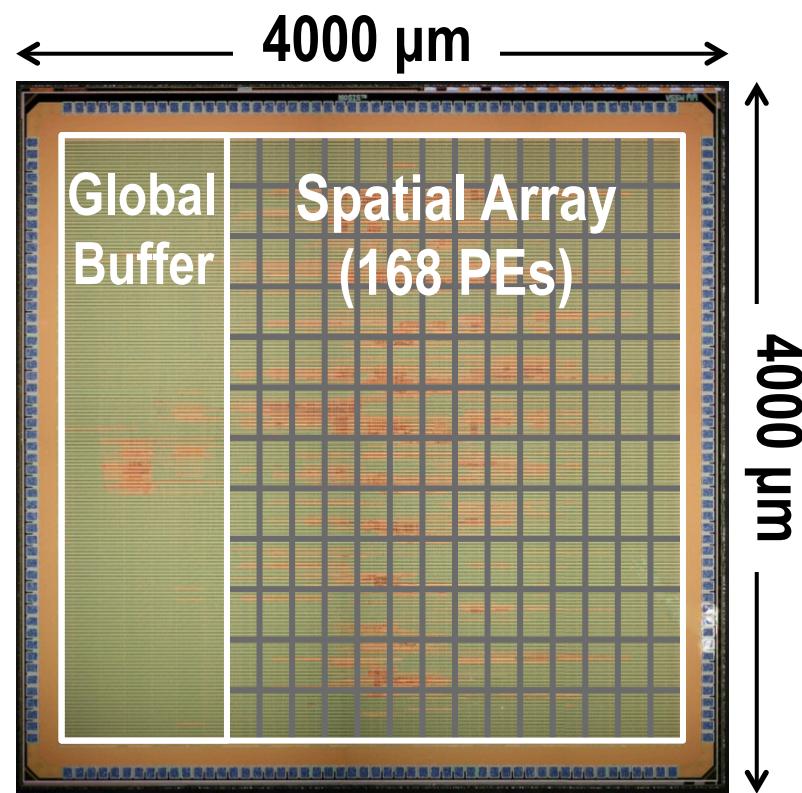
This talk will focus on the **Feature Extraction** cost

# Compare HOG vs. CNN

Compare using measured results from test chips (65 nm)



Object Detection using **HOG** features  
and Deformable Parts Models  
[VLSI 2016]



Eyeriss: **Convolutional Neural  
Networks**  
[ISSCC 2016, ISCA 2016]

# Hand-crafted Features (HOG)

HOG = Histogram of Oriented Gradients



Input Image

Gradient Vector

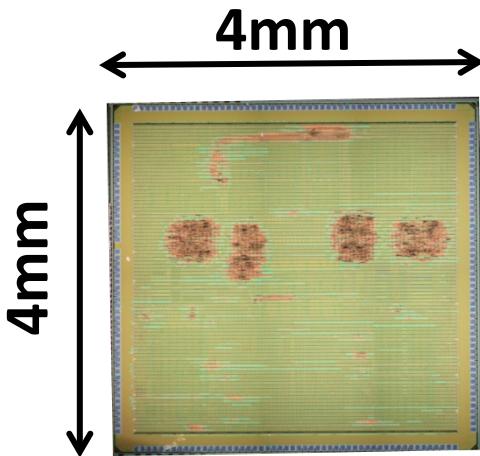


Cell Histogram

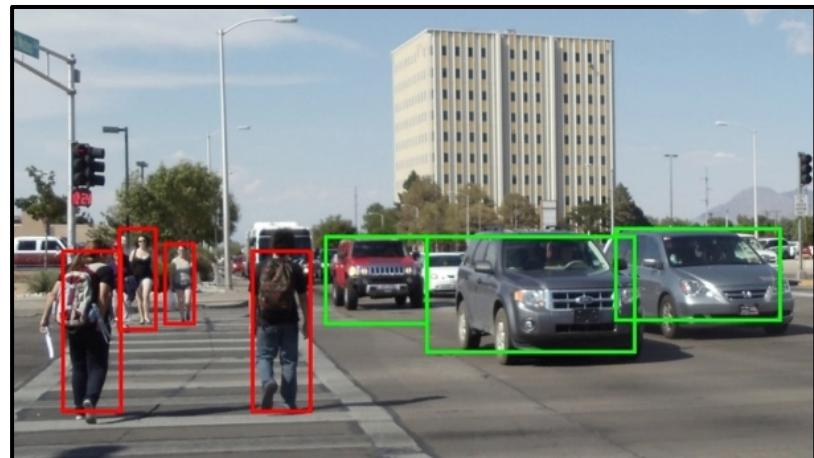


HOG Features

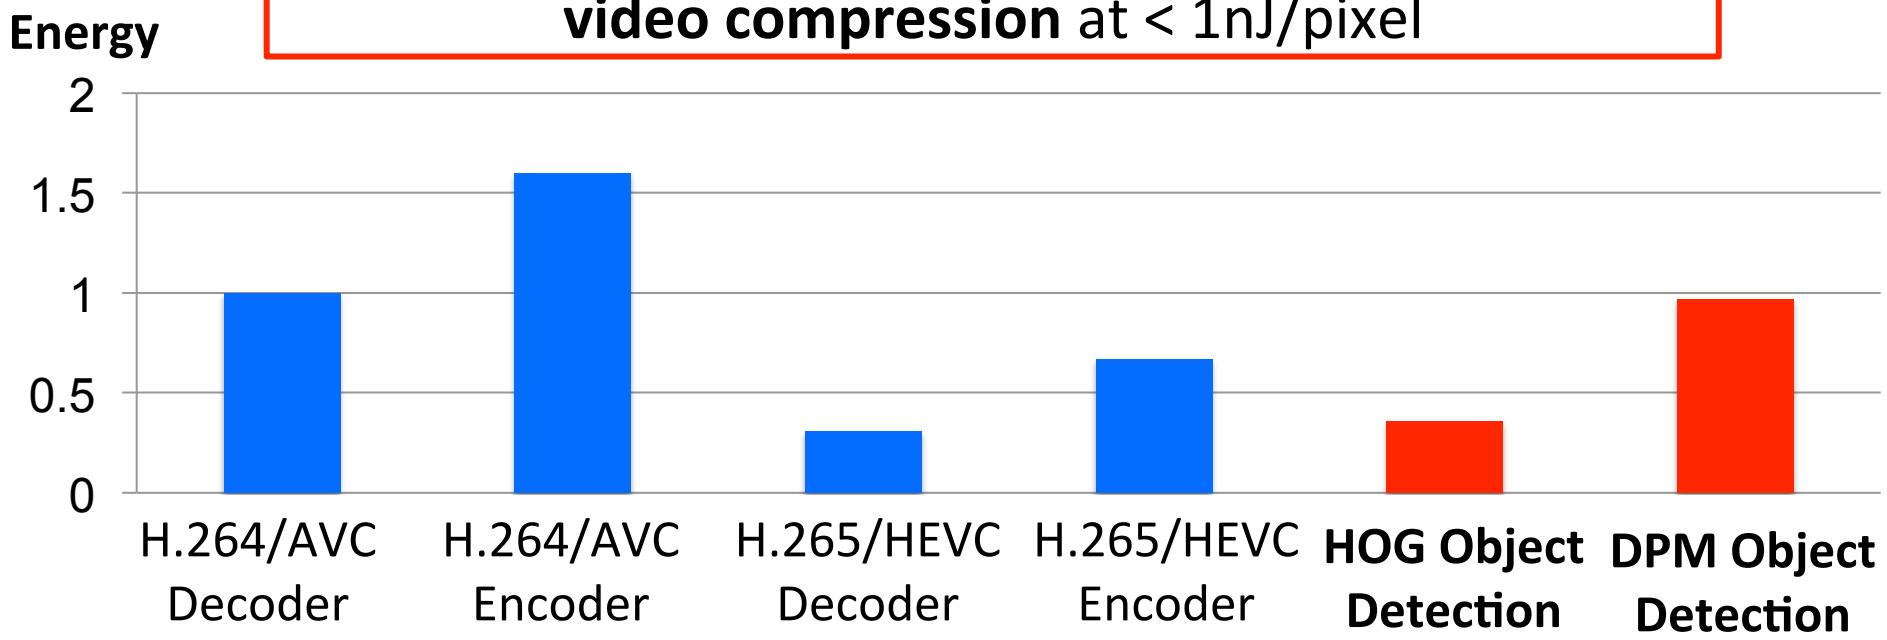
# Energy-Efficient Object Detection



MIT Object  
Detection Chip  
[VLSI 2016]

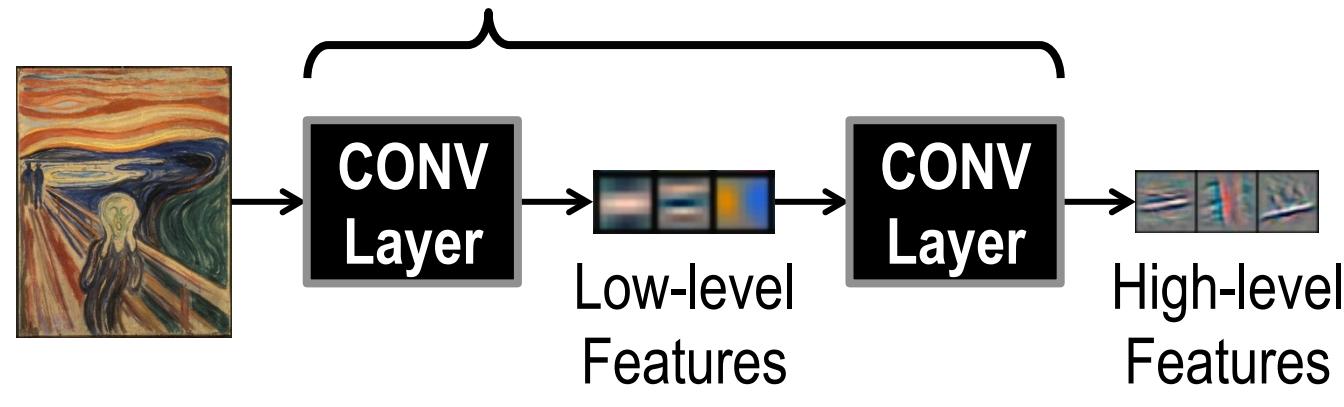


Enable object detection to be as **energy-efficient** as  
**video compression** at < 1nJ/pixel



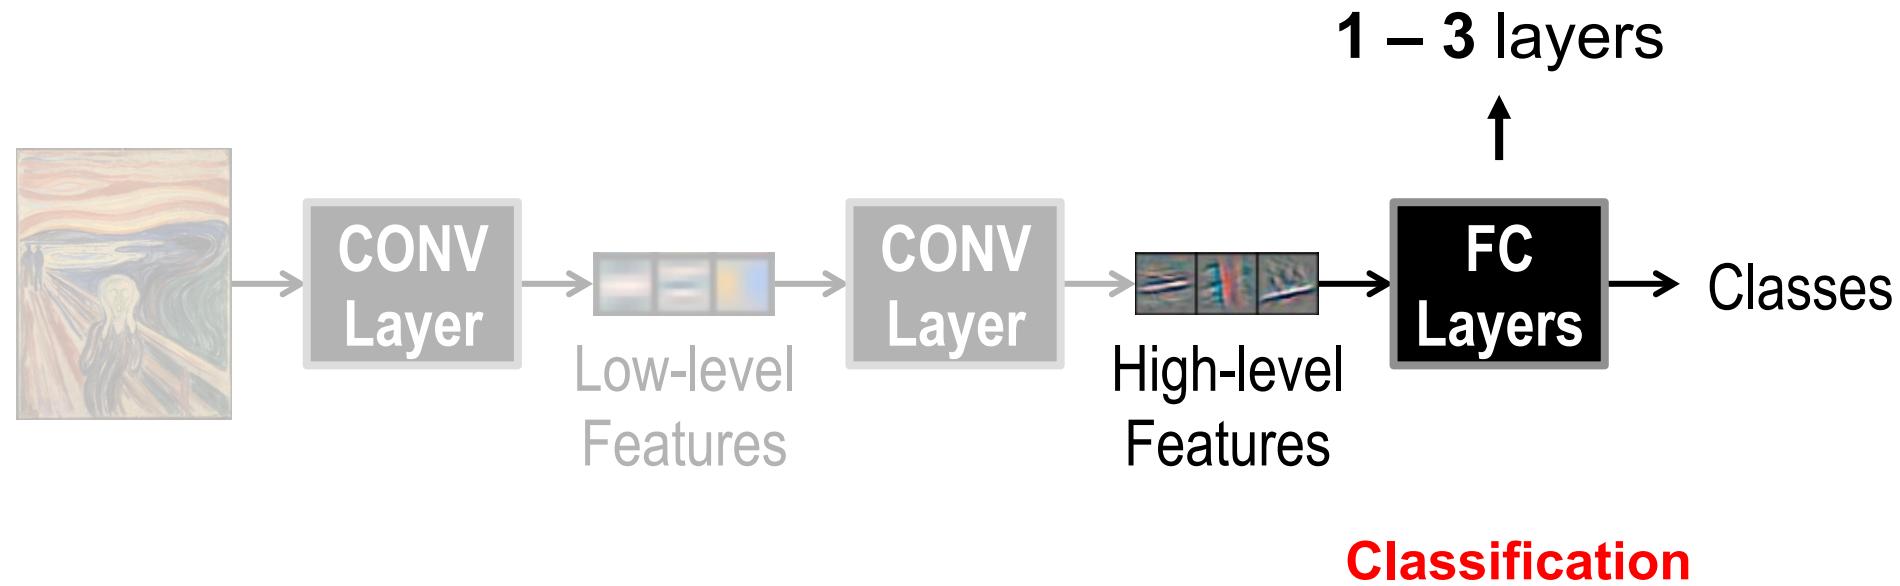
# Deep Convolutional Neural Networks

Modern *deep* CNN: up to 1000 CONV layers

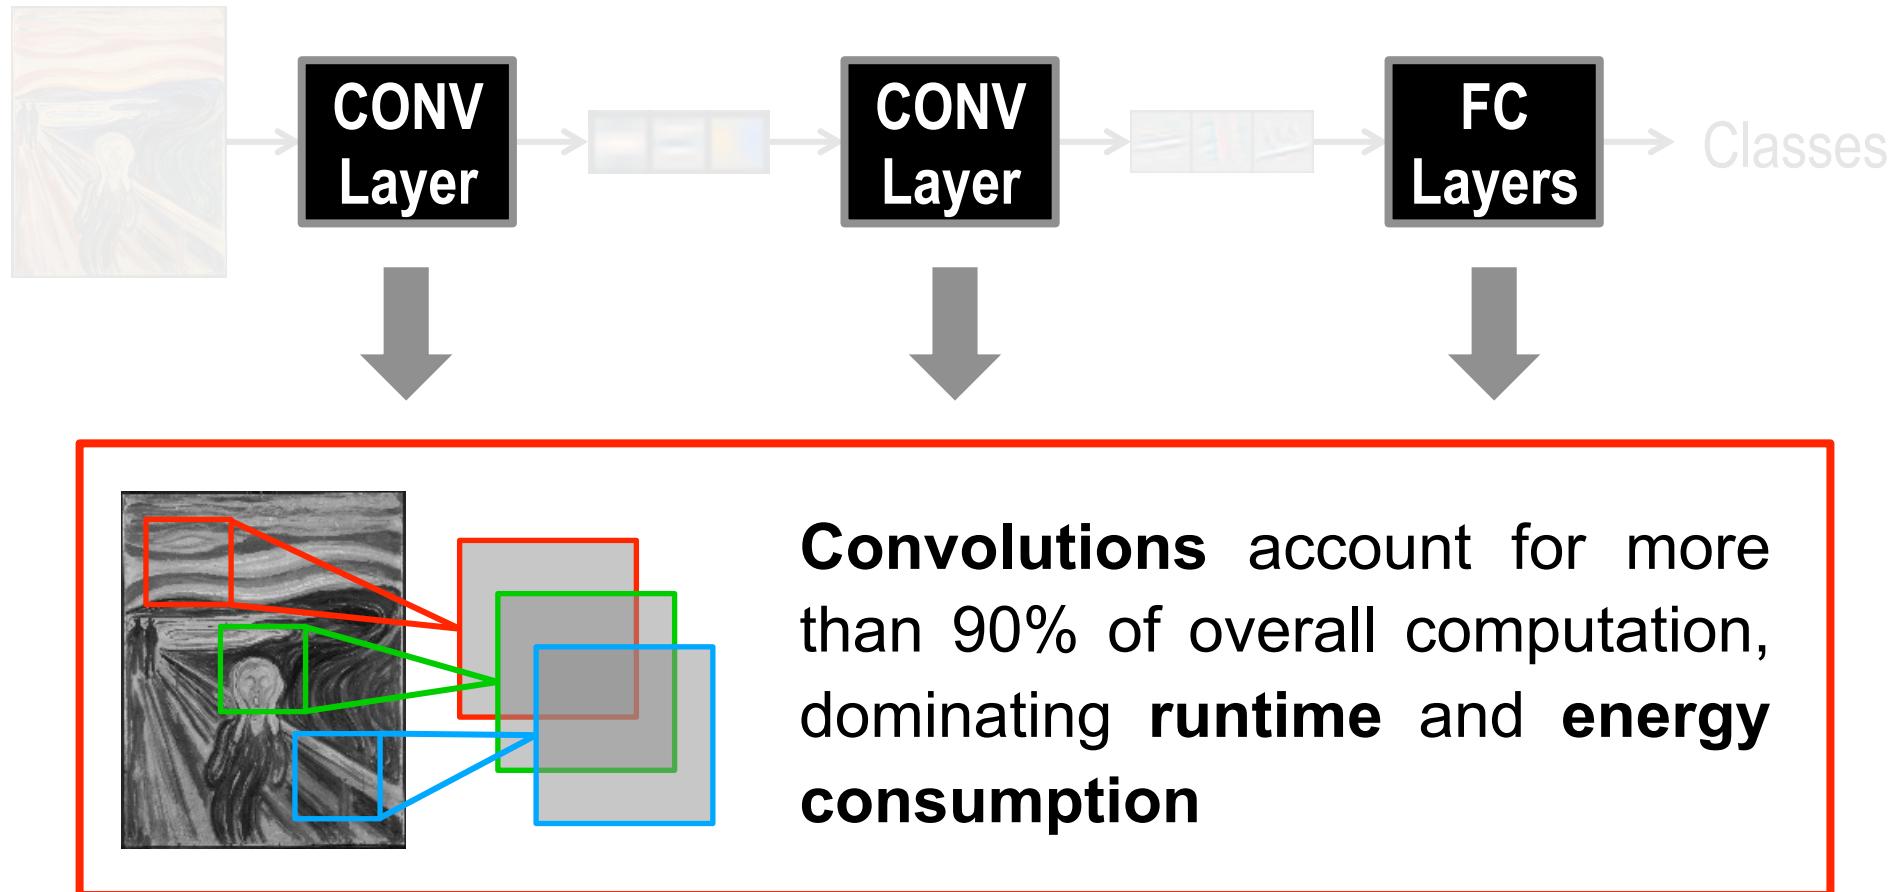


**Feature Extraction**

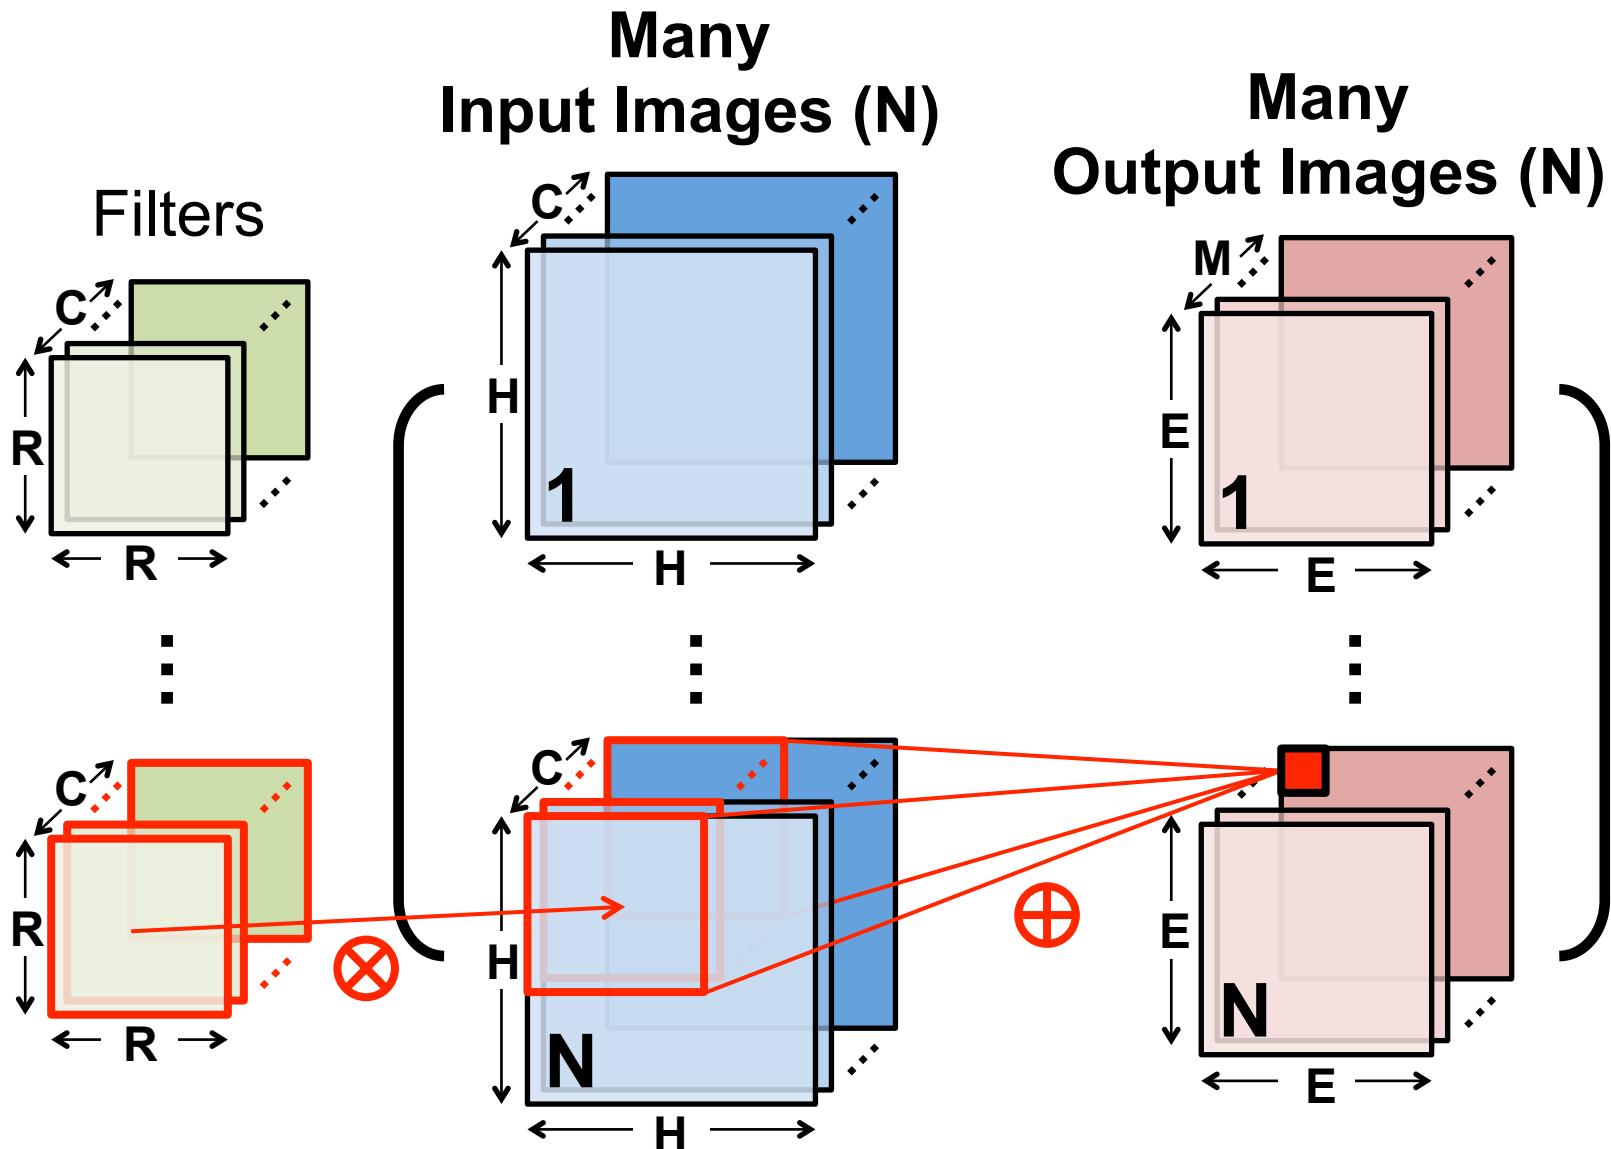
# Deep Convolutional Neural Networks



# Deep Convolutional Neural Networks



# High-Dimensional CNN Convolution

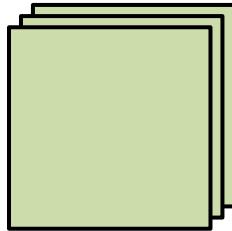


# Large Sizes with Varying Shapes

## AlexNet<sup>1</sup> Convolutional Layer Configurations

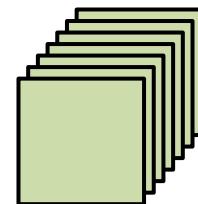
Layer	Filter Size (R)	# Filters (M)	# Channels (C)	Stride
1	11x11	96	3	4
2	5x5	256	48	1
3	3x3	384	256	1
4	3x3	384	192	1
5	3x3	256	192	1

Layer 1



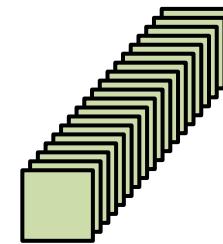
**34k Params**  
**105M MACs**

Layer 2



**307k Params**  
**224M MACs**

Layer 3



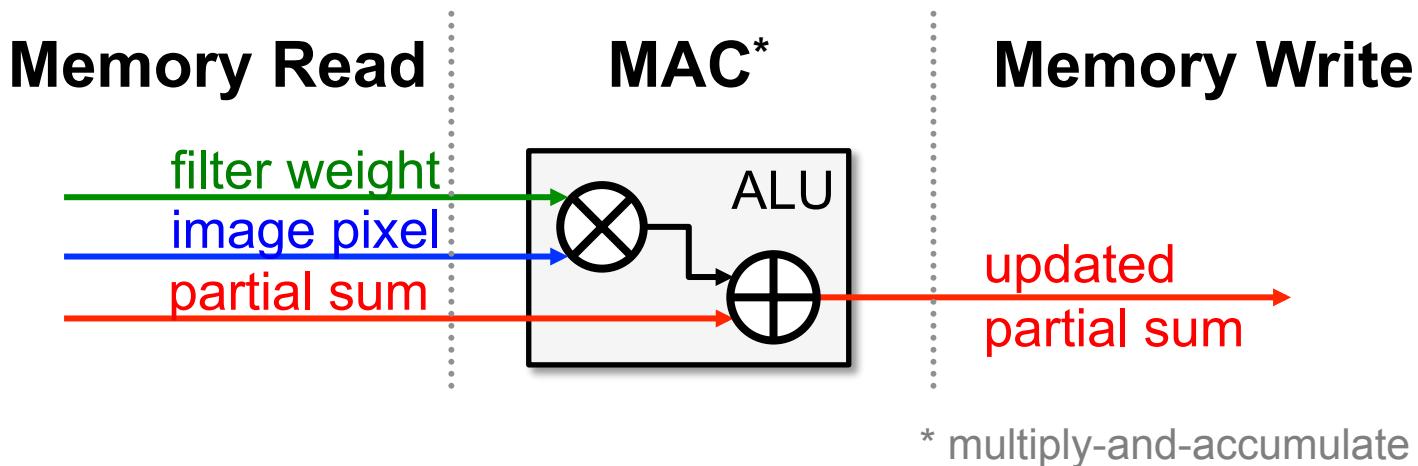
**885k Params**  
**150M MACs**

# Properties We Can Leverage

- Operations exhibit **high parallelism**  
→ **high throughput** possible

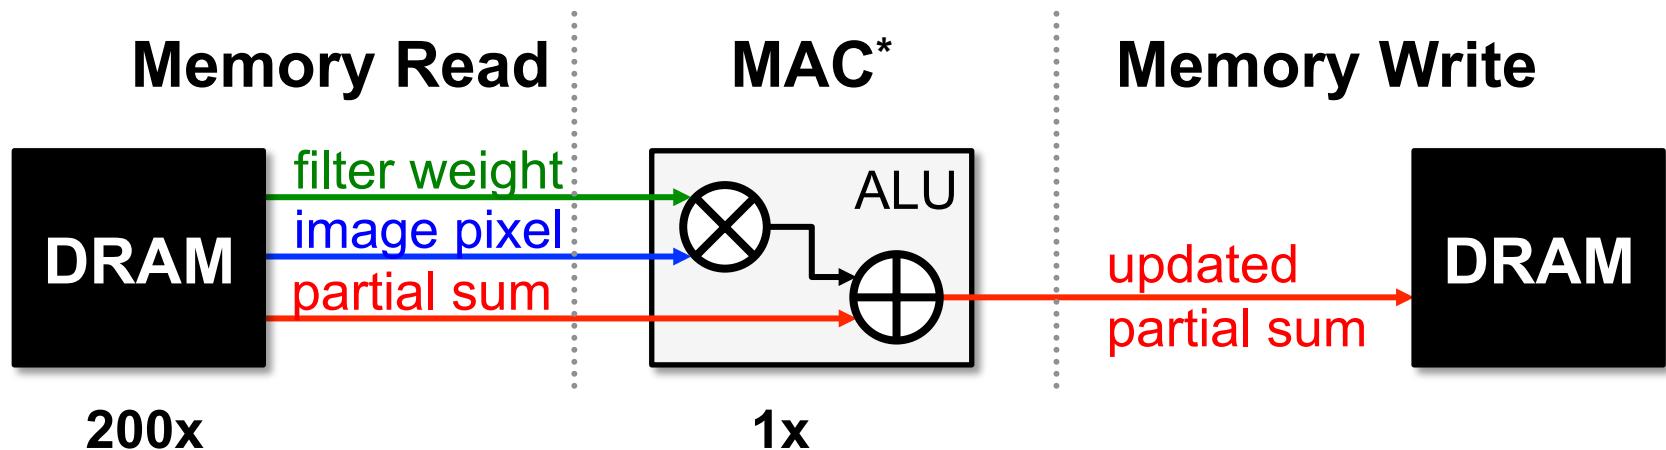
# Properties We Can Leverage

- Operations exhibit **high parallelism**  
→ **high throughput** possible
- Memory Access is the Bottleneck



# Properties We Can Leverage

- Operations exhibit **high parallelism**  
→ **high throughput** possible
- Memory Access is the Bottleneck

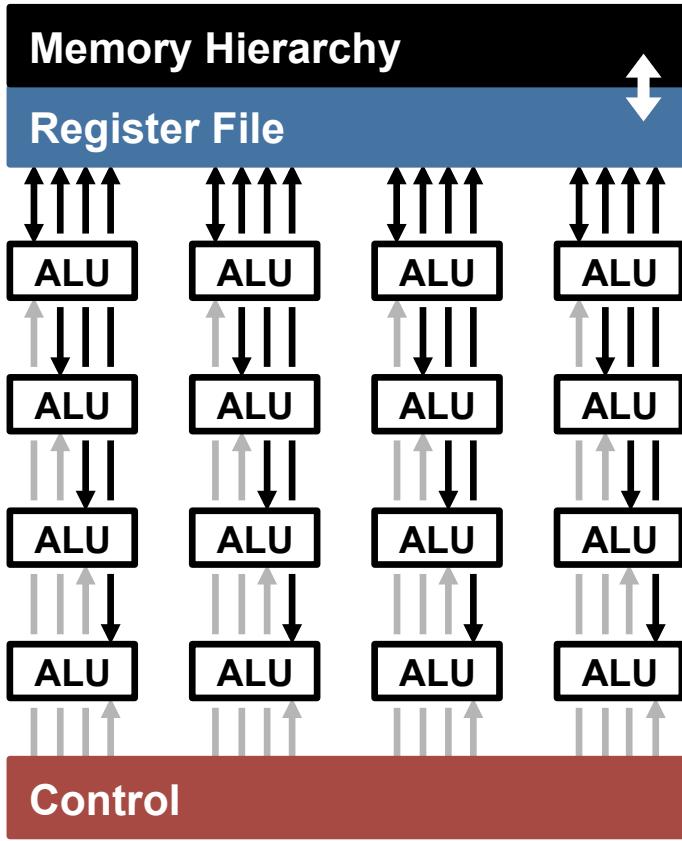


Worst Case: all memory R/W are **DRAM** accesses

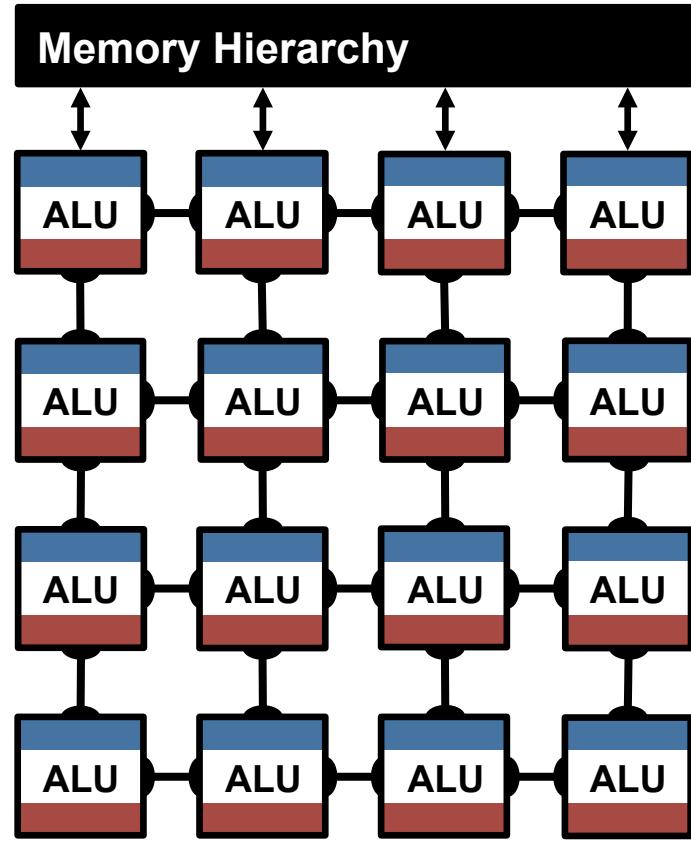
- Example: AlexNet [NIPS 2012] has **724M** MACs  
→ **2896M** DRAM accesses required

# Highly-Parallel Compute Paradigms

Temporal Architecture  
(SIMD/SIMT)



Spatial Architecture  
(Dataflow Processing)



# Advantages of Spatial Architecture

Temporal Architecture  
(SIMD/SIMT)

## Efficient Data Reuse

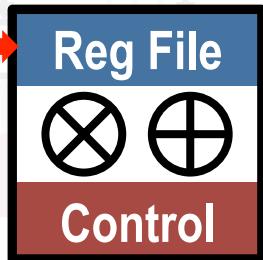
Distributed local storage (RF)

## Inter-PE Communication

Sharing among regions of PEs

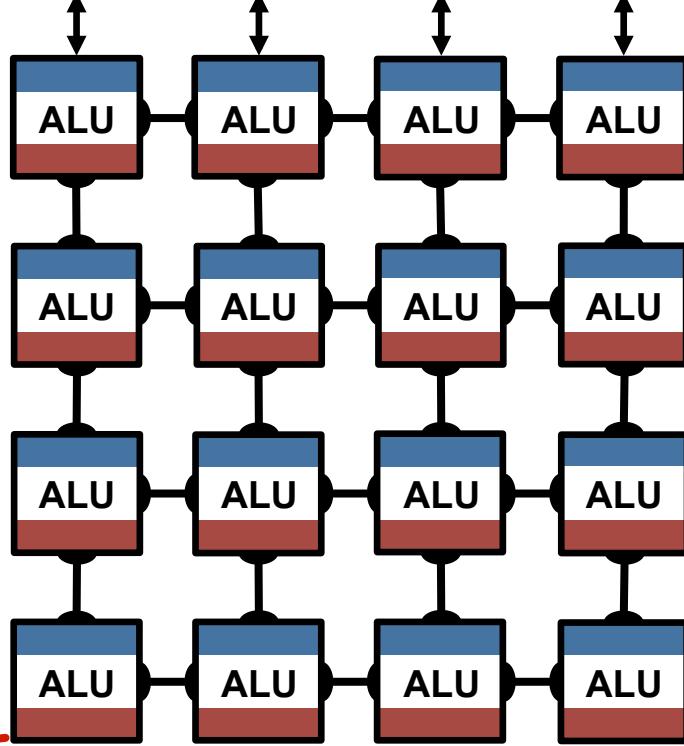
## Processing Element (PE)

0.5 – 1.0 kB

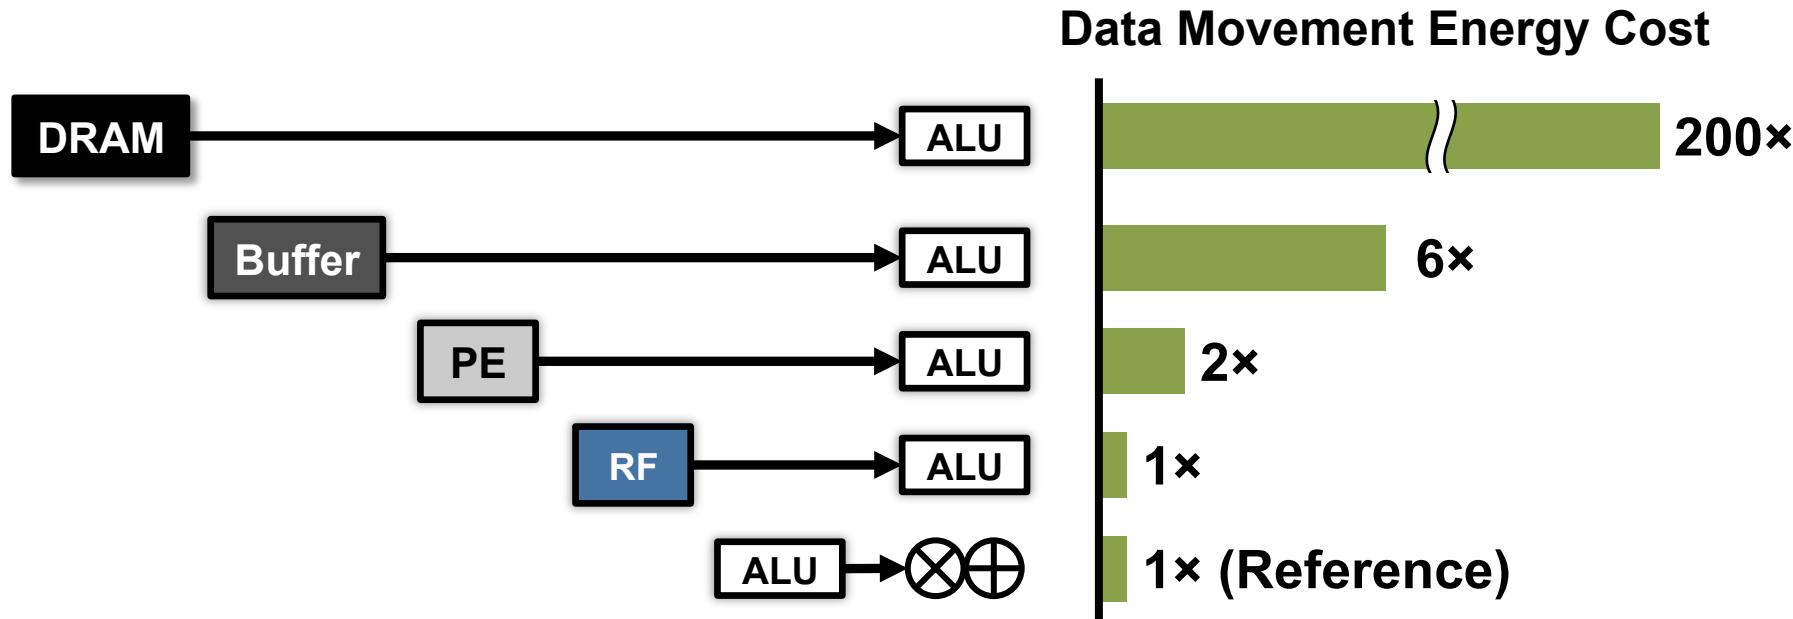
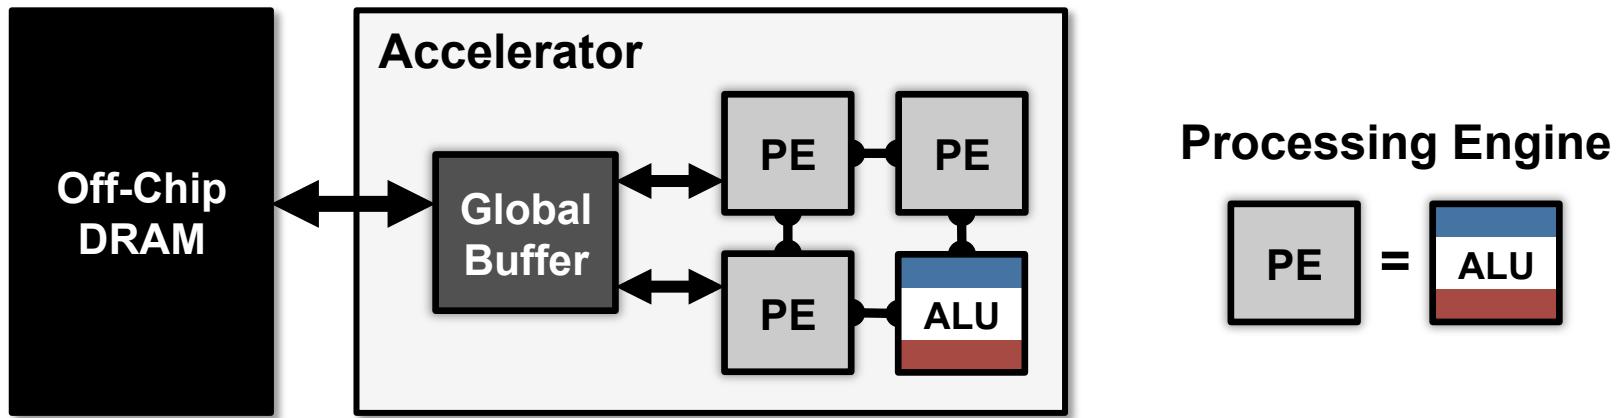


Spatial Architecture  
(Dataflow Processing)

## Memory Hierarchy



# Data Movement is Expensive

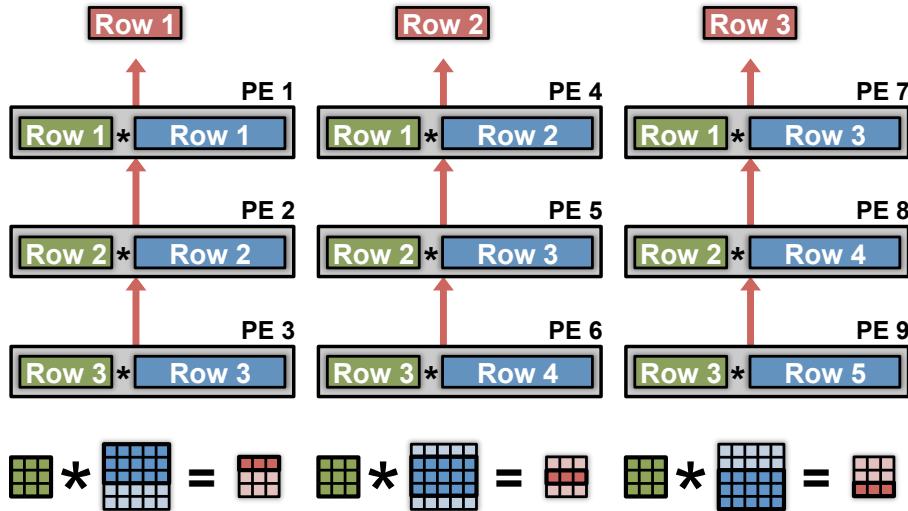


Maximize data reuse at lower levels of hierarchy

# Optimization to Reduce Data Movement

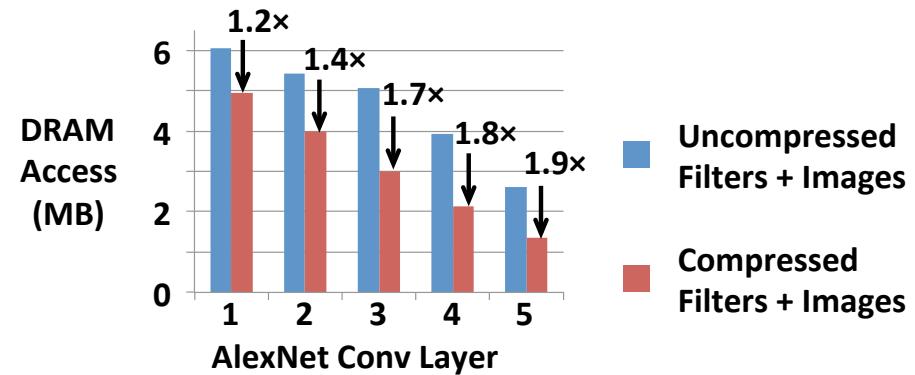
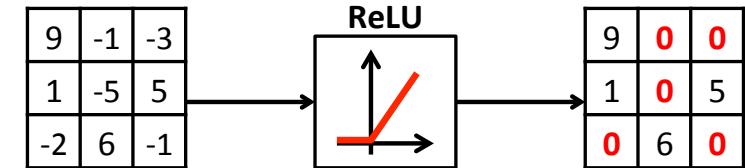
- Energy-efficient dataflow to reduce data movement
- Exploit data statistics for high energy efficiency

## Row Stationary Dataflow

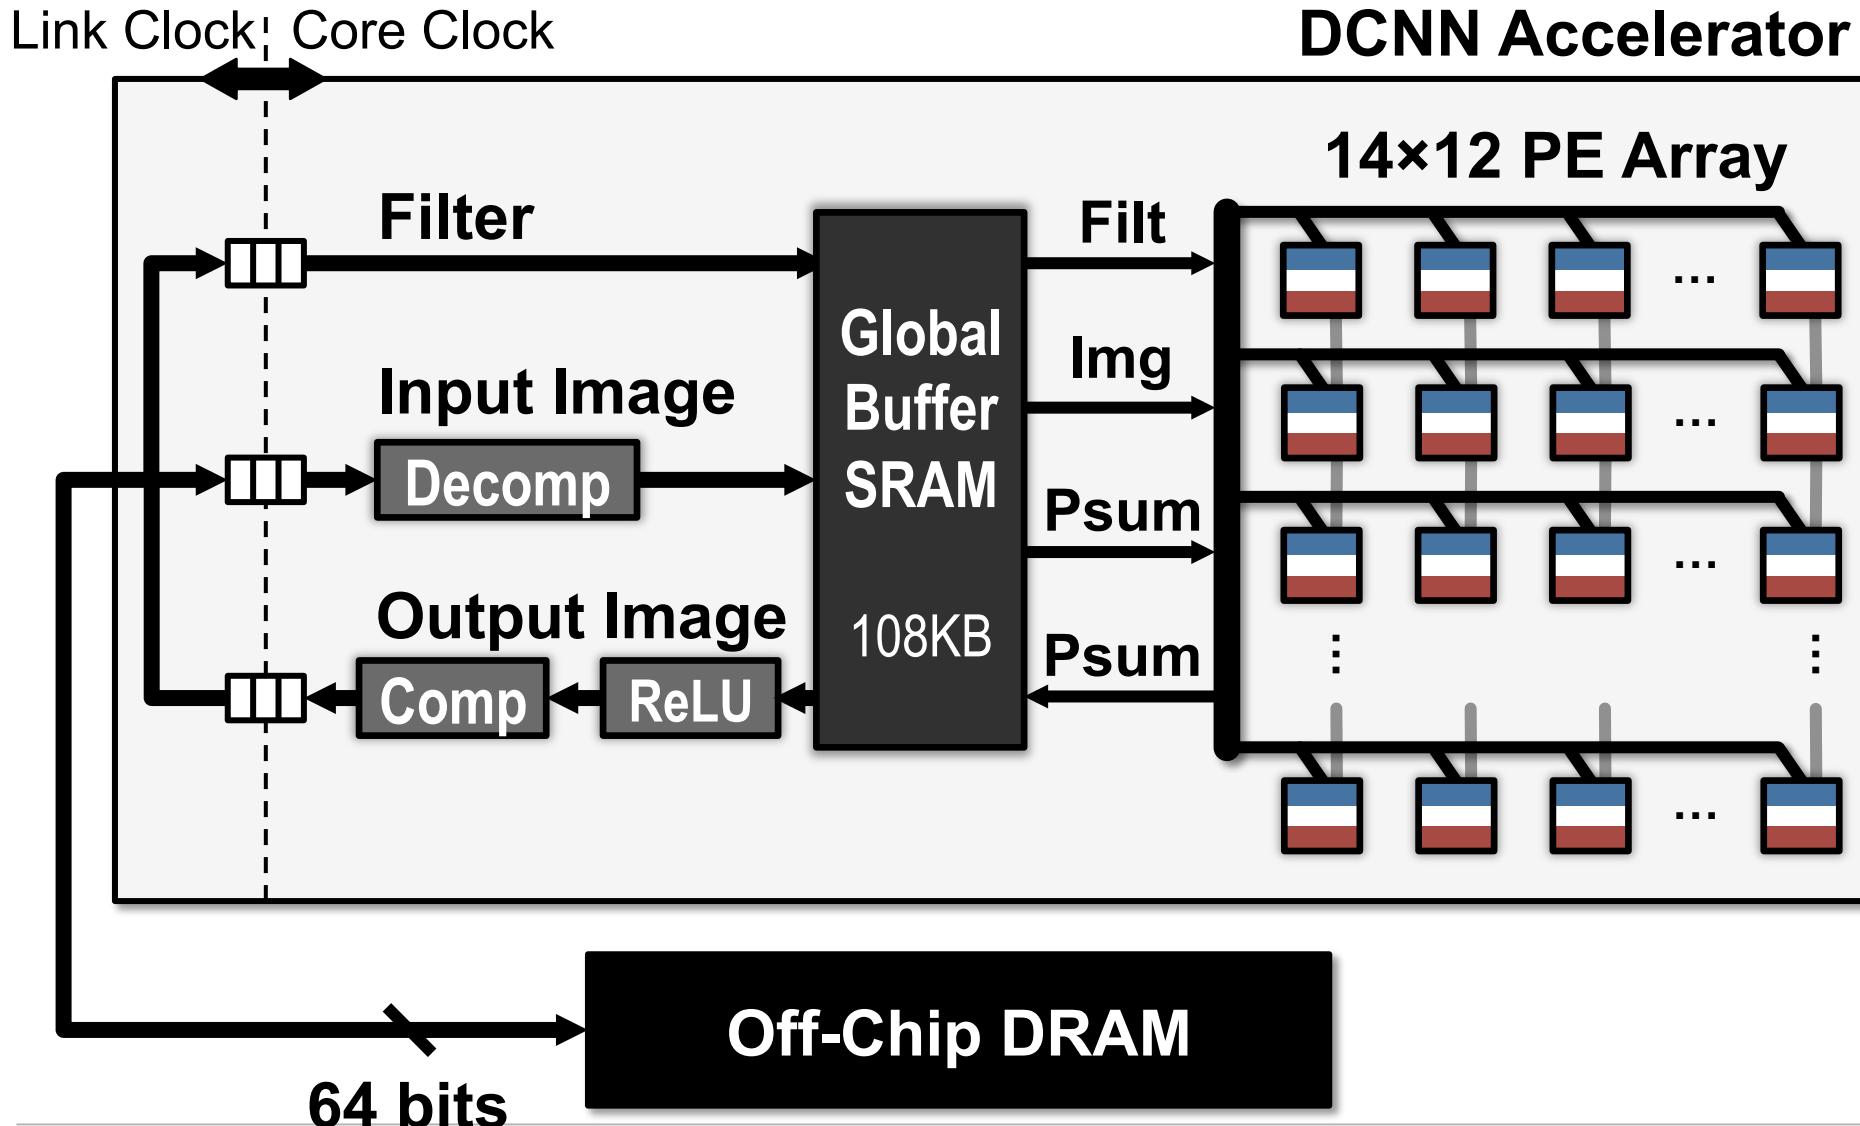


## Sparsity in Activations

Apply Non-Linearity (ReLU) on Filtered Image Data

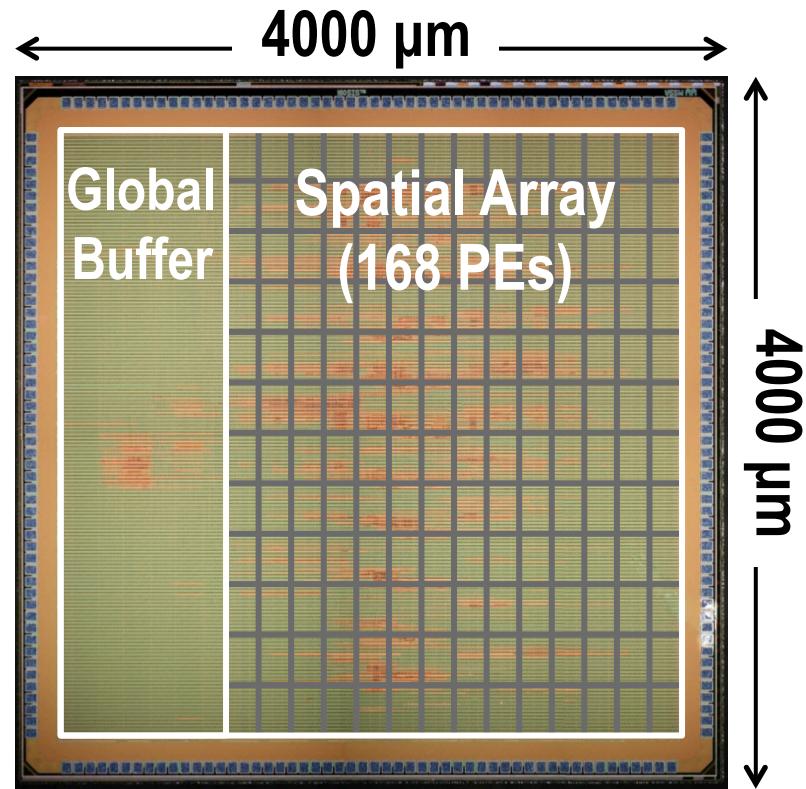


# Eyeriss Deep CNN Accelerator



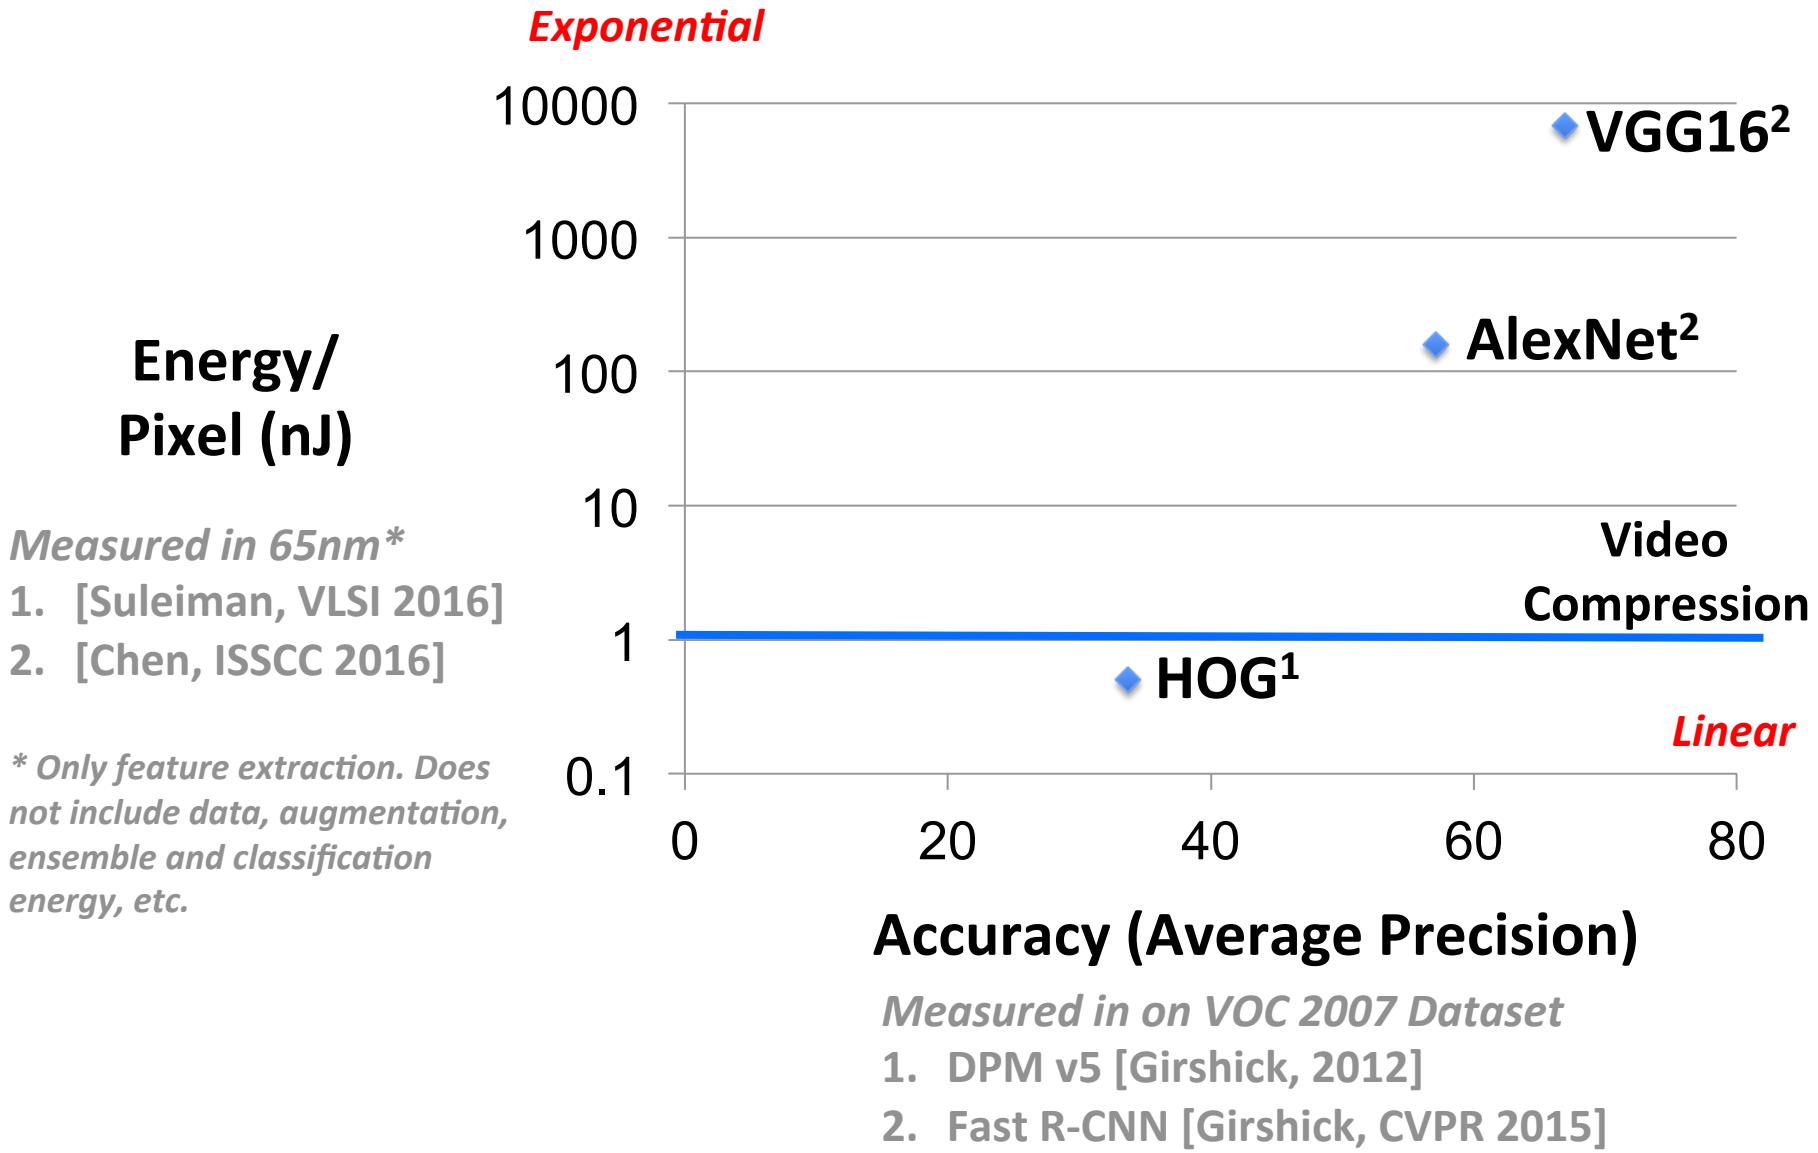
# Eyeriss Chip Spec & Measurement Results

<b>Technology</b>	TSMC 65nm LP 1P9M
<b>On-Chip Buffer</b>	108 KB
<b># of PEs</b>	168
<b>Scratch Pad / PE</b>	0.5 KB
<b>Core Frequency</b>	100 – 250 MHz
<b>Peak Performance</b>	33.6 – 84.0 GOPS
<b>Word Bit-width</b>	16-bit Fixed-Point
<b>Natively Supported CNN Shapes</b>	Filter Width: 1 – 32 Filter Height: 1 – 12 Num. Filters: 1 – 1024 Num. Channels: 1 – 1024 Horz. Stride: 1–12 Vert. Stride: 1, 2, 4

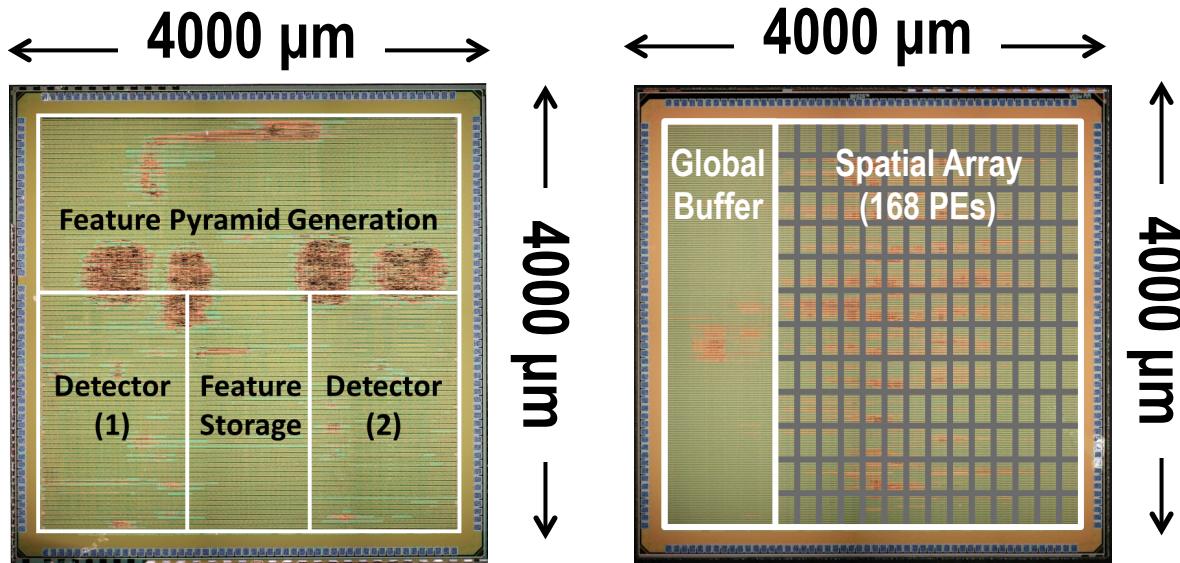


Over 10x more energy efficient than a mobile GPU (Nvidia TK1)

# Features: Energy vs. Accuracy



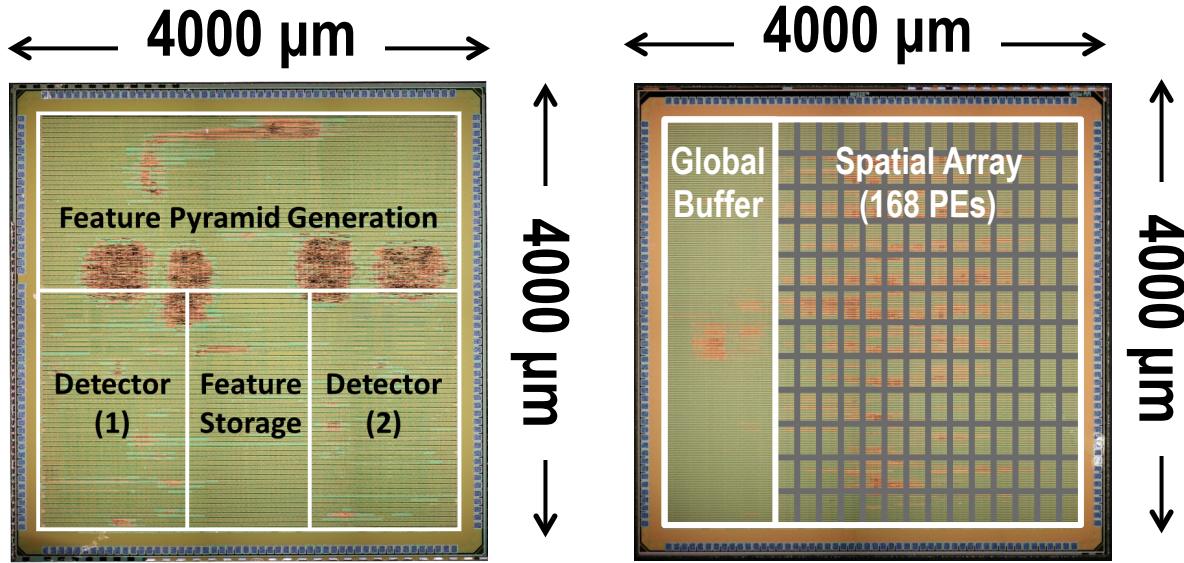
# HOG vs. CNN: Hardware Cost



	<b>HOG [VLSI 2016]</b>	<b>CNN [ISSCC 2016]</b>
<b>Technology</b>	TSMC LP 65nm	TSMC LP 65m
<b>Gate Count (kgates)</b>	893	1176
<b>Memory (kB)</b>	159	181.5

Similar Hardware Cost (comparable with Video Compression)

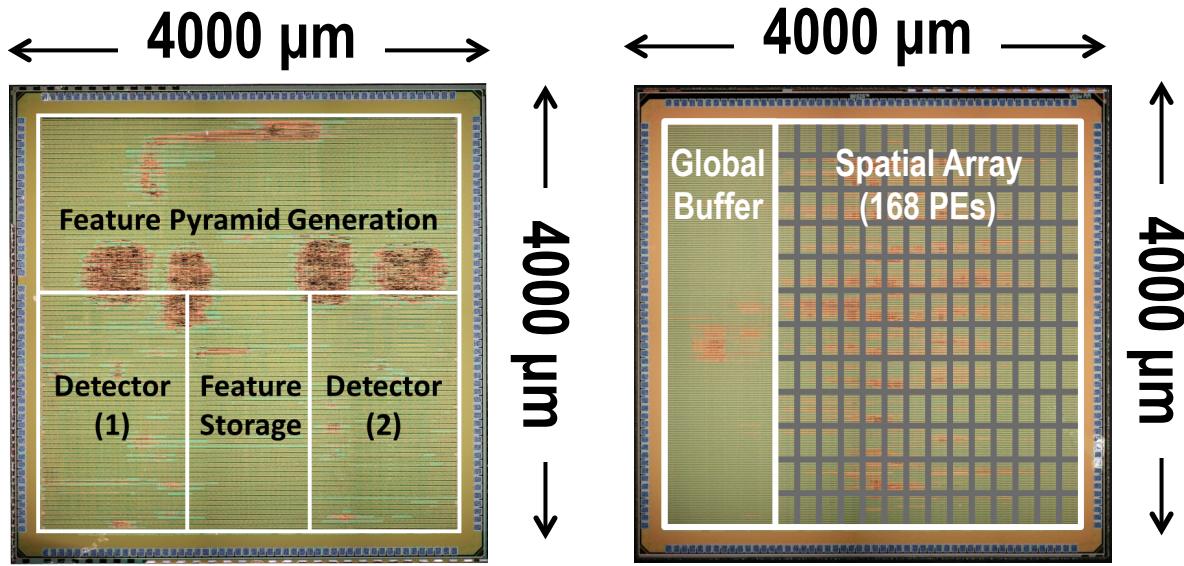
# HOG vs. CNN: Throughput



	<b>HOG</b>	<b>CNN (AlexNet)</b>	<b>CNN (VGG-16)</b>
<b>Throughput (Mpixels/s)</b>	62.5	1.8	0.04
<b>GOP/Mpixel</b>	0.7	25.8	610.3
<b>Throughput (GOPS)</b>	46.0	46.2	21.4

Throughput gap explained by GOP/Mpixel gap

# HOG vs. CNN: Energy and DRAM Access



	<b>HOG</b>	<b>CNN (AlexNet)</b>	<b>CNN (VGG-16)</b>
<b>Energy (nJ/pixel)</b>	0.5	155.5	6742.9
<b>GOP/Mpixel</b>	0.7	25.8	610.3
<b>Energy (GOPs/W)</b>	1570	166.2	90.7
<b>DRAM (B/pixel)</b>	1.0	74.7	2128.6

Energy gap larger than GOPS/W gap

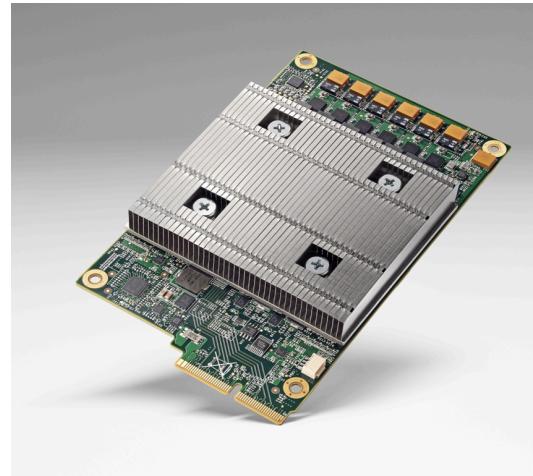
# Energy Gap between CNN and HOG

- **CNNs require more operations per pixel**
  - AlexNet vs. HOG = **37x**
  - VGG-16 vs. HOG = **872x**
- **CNN requires a programmable architecture**
  - **Example:** AlexNet CONV layers have **2.3M weights** (assume 8-bits per weight); Area budget of HOG chip is ~**1000 kgates, 150kB**
  - **Design A: Hard-wired weights**
    - Only have 10k multipliers with fixed weights (**>100x increase in area**)
  - **Design B: Store all weights on-chip**
    - Only store 150k weights on chip (**>10x increase in storage**)
  - **Support different shapes per layer and different weights**

# Closing the Energy Gap

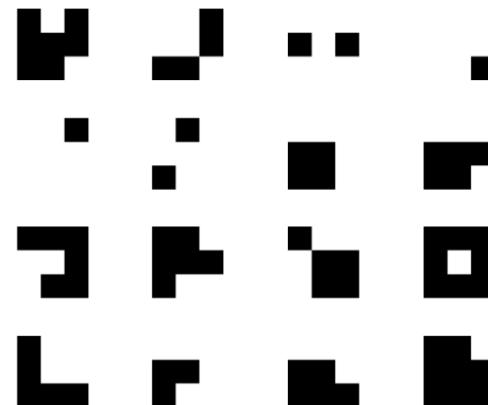
# Methods to Reduce Energy of CNNs

- **Reduce Precision**
  - [Google TPU, ISCA 2017], [XNOR-Net, ECCV 2016], [BinaryNets, arXiv 2016]
- **Sparsity by Pruning**
- **Data Compression**
  - [Chen, ISSCC 2016], [Han, ISCA 2016], [Moons, VLSI 2016]
- **Energy Optimized Dataflow**
  - [Chen, ISCA 2016]



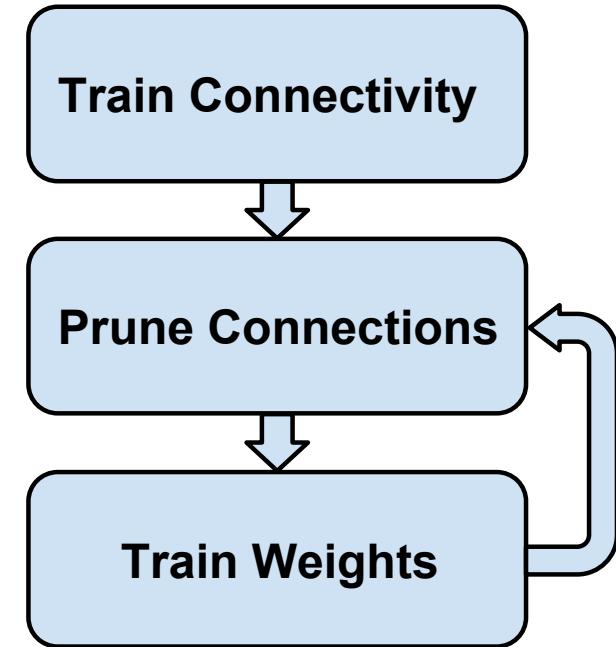
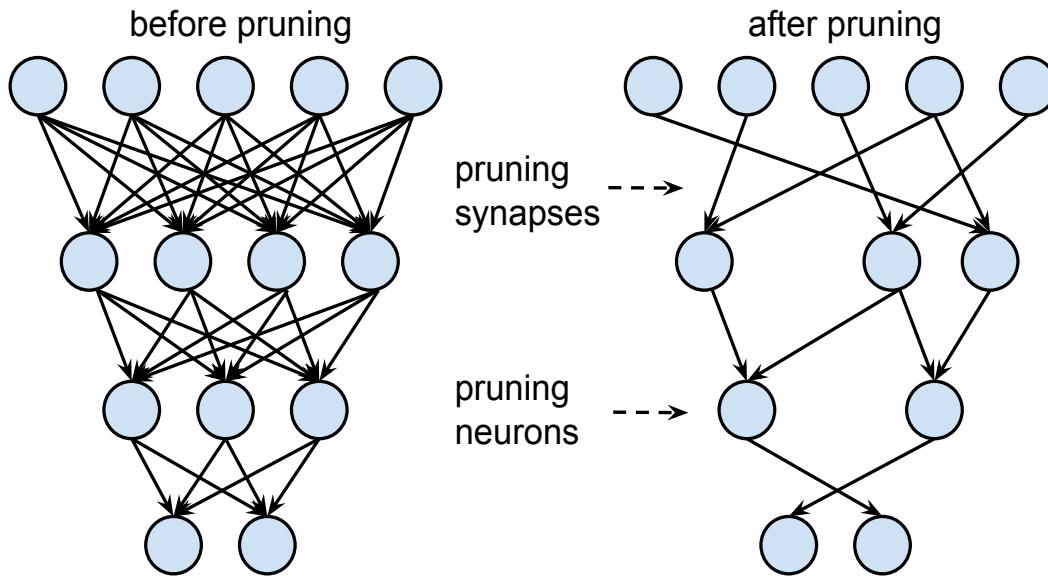
**Google's TPU (8-bits)**

*Binary Filters*



# Pruning – Make Weights Sparse

Prune based on *magnitude* of weights



**Example:** AlexNet

**Weight Reduction:** CONV layers 2.7x, FC layers 9.9x

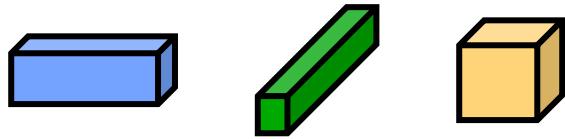
(*Most reduction on fully connected layers*)

**Overall:** 9x weight reduction, 3x MAC reduction

# Key Metrics for Embedded DNN

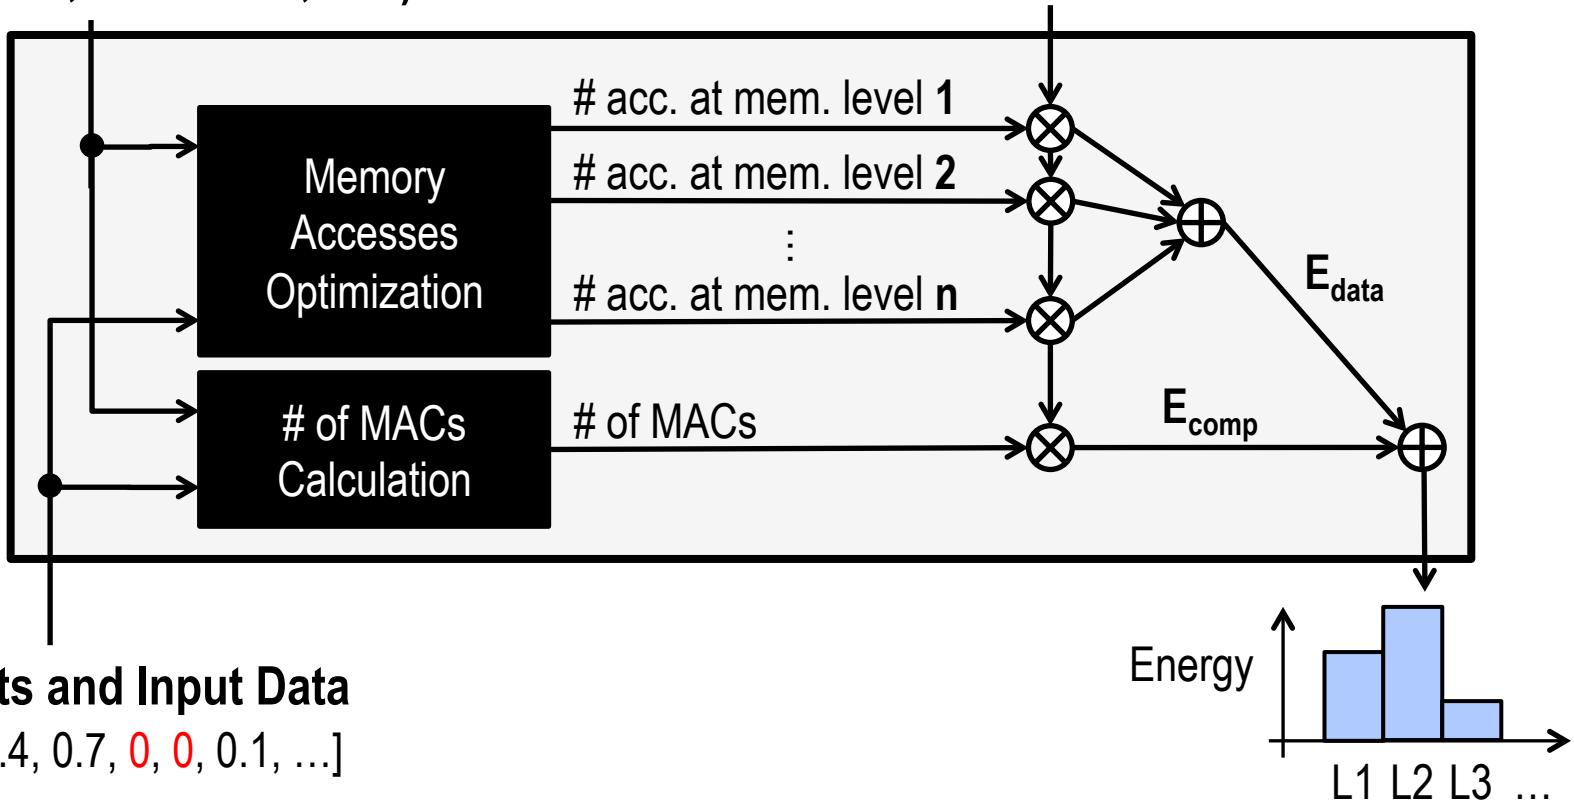
- Accuracy → Measured on Dataset
- Speed → Number of MACs
- Storage Footprint → Number of Weights
- Energy → ?

# Energy-Evaluation Methodology



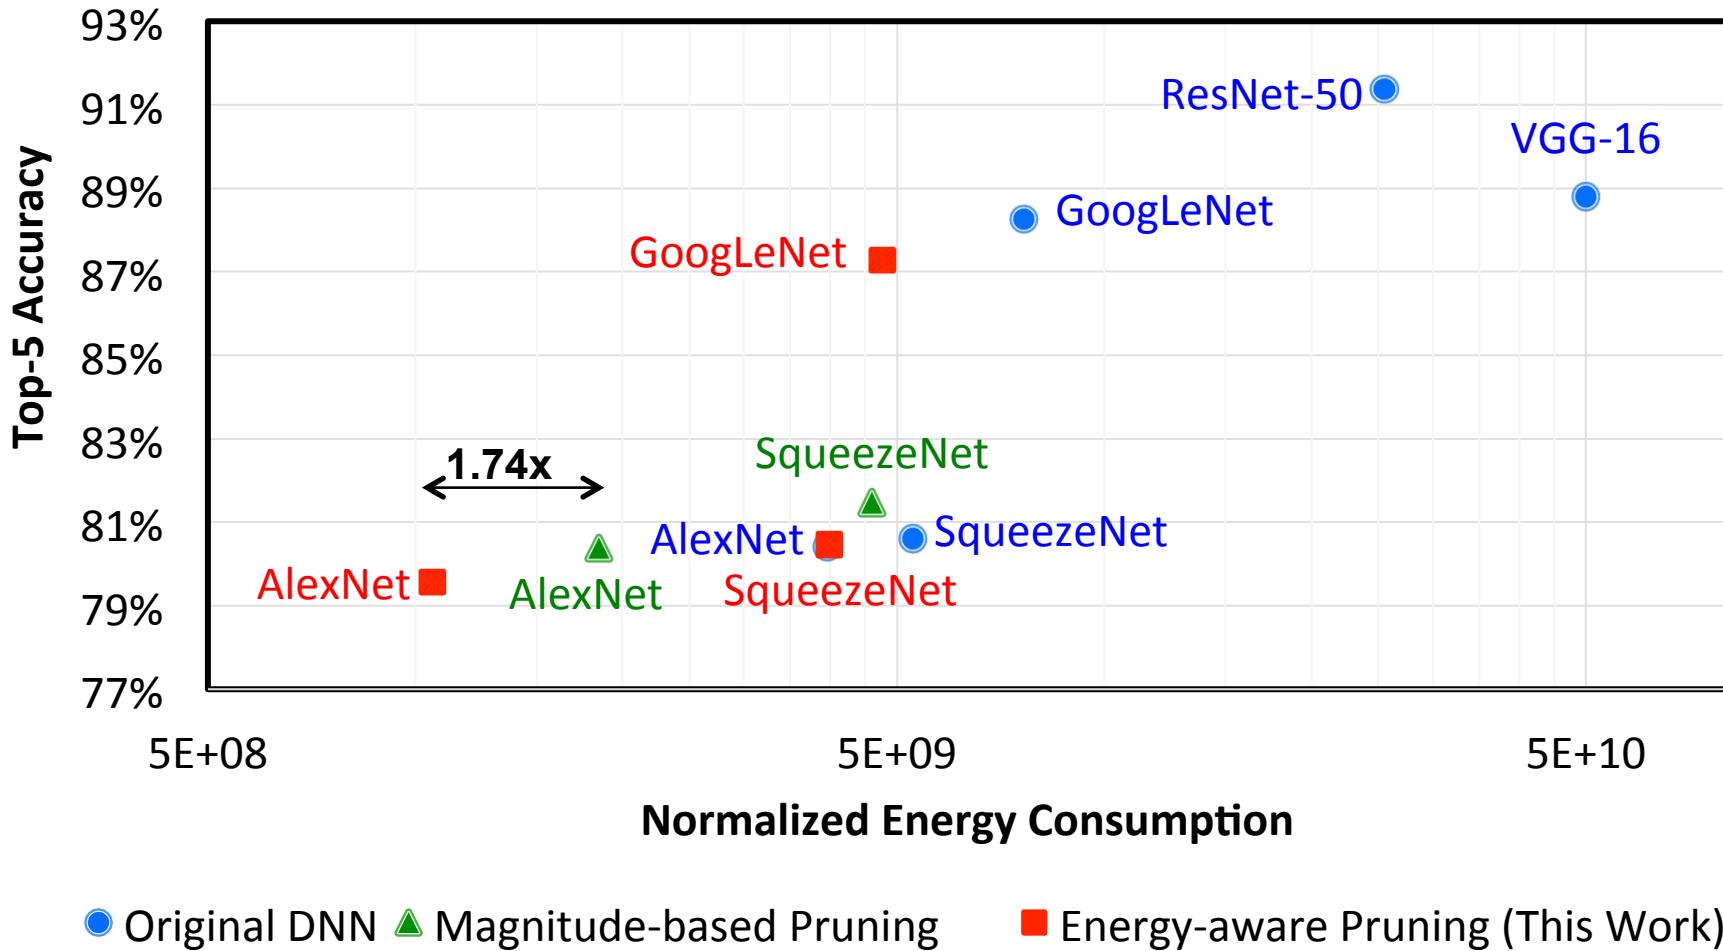
**CNN Shape Configuration**  
(# of channels, # of filters, etc.)

**Hardware Energy Costs of each MAC and Memory Access**



[Yang et al., CVPR 2017]

# Energy-Aware Pruning



Remove weights from layers in order of highest to lowest energy  
**3.7x reduction in AlexNet / 1.6x reduction in GoogLeNet**

# Summary

- CNN gives higher accuracy than HOG features (2x) at the cost of increase energy (311x to 13486x)
- Energy gap due to (1) CNN requires more operations per pixel and (2) CNN requires a programmable architecture
- Joint algorithm and hardware design can deliver additional energy savings to help close this gap

More info about **Eyeriss** and **Tutorial on DNN Architectures** at <http://eyeriss.mit.edu>



V. Sze, Y.-H. Chen, T-J. Yang, J. Emer, “*Efficient Processing of Deep Neural Networks: A Tutorial and Survey*”, arXiv, 2017

 Follow @eems\_mit